

# 9. Generative Classifiers

## STA3142 Statistical Machine Learning

**Kibok Lee**

Assistant Professor of

Applied Statistics / Statistics and Data Science

Mar {26, 28}, 2024




**연세대학교**  
YONSEI UNIVERSITY

# Assignment 1


- Due **Friday 3/29, 11:59pm**
- Topics
  - (Programming) NumPy basics
  - (Programming) Linear regression on a polynomial
  - (Math) Derivation and proof for linear regression
- Please read the instruction carefully!
  - Submit one pdf and one zip file separately
  - Write your code only in the designated spaces
  - Do not import additional libraries
  - ...
- If you feel difficult, consider to take **option 2**.

# Recap: Supervised Learning

- Learning a function  $h: \mathcal{X} \rightarrow \mathcal{Y}$
- Labels could be discrete or continuous
  - Discrete labels: **classification**
  - Continuous labels: **regression**



classification



regression

# Classification Strategies

- Learning the distributions  $p(C_k|x)$ 
  - Discriminative models: Directly model  $p(C_k|x)$  and learn parameters from the training set.
  - Generative models: Learn class densities  $p(x|C_k)$  and priors  $p(C_k)$  to obtain  $p(x, C_k) = p(x|C_k)p(C_k)$
- Nearest neighbor classification
  - Given query data  $x$ , find the closest training points and do majority vote.
- Discriminant functions
  - Learn a function  $h(x)$  that maps  $x$  onto some  $C_k$ .

# Outline

- Generative models: Learn class densities  $p(x|C_k)$  and priors  $p(C_k)$  s.t.  $p(x, C_k) = p(x|C_k)p(C_k)$ 
  - Gaussian Discriminant Analysis
  - Naïve Bayes Classifier

# Probabilistic Generative Models

- Bayes' theorem reduces the classification problem  $p(C_k|x)$  to estimating the distribution of the data.
- Density estimation problems are easy to learn from labeled training data.
  - Priors:  $p(C_k)$
  - Class densities:  $p(x|C_k)$
- Learning: Maximum likelihood estimation (MLE)
- Classification: Maximum a posteriori (MAP) estimation

$$\operatorname{argmax}_k p(C_k|x) = \operatorname{argmax}_k p(C_k, x)$$

# Probabilistic Generative Models

- For two-class classification, Bayes' theorem says:

$$p(C_1|\mathbf{x}) = \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_1)p(C_1) + p(\mathbf{x}|C_2)p(C_2)}$$

- The posterior is then expressed as the sigmoid of log odds:

$$p(C_1|\mathbf{x}) = \frac{1}{1 + \exp(-a)} = \sigma(a)$$

$$\text{where } a = \ln \frac{p(C_1|\mathbf{x})}{p(C_2|\mathbf{x})} = \ln \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_2)p(C_2)}$$

# Generative vs. Discriminative

- The **generative** approach is typically model-based and makes it possible to generate synthetic data from  $p(x|C_k)$ .
  - By comparing the synthetic data and real data, we get a sense of how good the generative model is.
- The **discriminative** approach typically has fewer parameters to estimate and have less assumptions about the data distribution (i.e., no  $p(x, \dots)$ ).
  - Linear (e.g., logistic regression) or quadratic (e.g., Gaussian discriminant analysis) in the input.
  - Less generative assumptions about the data (i.e., constructing the features may need prior knowledge)



# Gaussian Discriminant Analysis

# Gaussian Discriminant Analysis

- Prior distribution  $p(C_k)$ : Constant (e.g., Bernoulli)
- Likelihood  $p(\mathbf{x}|C_k)$ : Gaussian distribution

$$p(\mathbf{x}|C_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2}(\mathbf{x} - \mu_k)^T \Sigma^{-1}(\mathbf{x} - \mu_k) \right\}$$

- Classification using Bayes' rule:




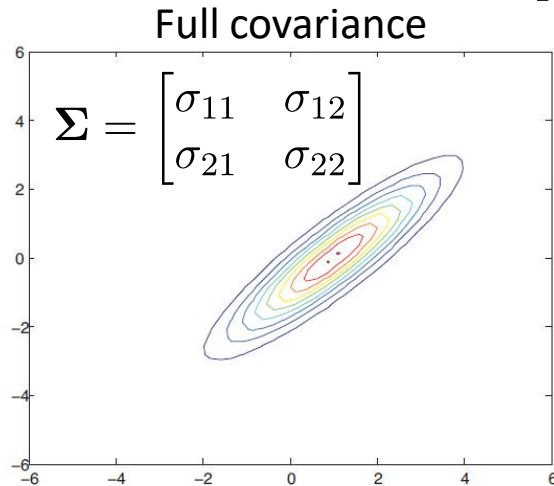
$$p(C_k|\mathbf{x}) = p(\mathbf{x}|C_k)p(C_k)/p(\mathbf{x})$$

- For two-class classification,  $p(C_1|\mathbf{x}) = \sigma(a)$

$$\text{where } a = \ln \frac{p(C_1|\mathbf{x})}{p(C_2|\mathbf{x})} = \ln \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_2)p(C_2)}$$


# Examples of Gaussian Distributions

- Probability density  $p(\mathbf{x})$  for 2-dim case



# Example: Two-Class GDA

- GDA assumes the **same covariance** for all classes.
  - The decision boundary is linear.
  - e.g., two-class classification



# Two-Class GDA Formulation

- We model  $p(\mathbf{x}|C_k)$  as Gaussian distributions with the **same covariance** matrix.

$$p(\mathbf{x}|C_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\boldsymbol{\Sigma}|^{1/2}} \exp \left\{ -\frac{1}{2}(\mathbf{x} - \mu_k)^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \mu_k) \right\}$$

- Then, the posterior is derived as

$$p(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

where  $\mathbf{w} = \boldsymbol{\Sigma}^{-1}(\mu_1 - \mu_2)$

$$w_0 = -\frac{1}{2}\mu_1^T \boldsymbol{\Sigma}^{-1} \mu_1 + \frac{1}{2}\mu_2^T \boldsymbol{\Sigma}^{-1} \mu_2 + \ln \frac{p(C_1)}{p(C_2)}$$

# Two-Class GDA Derivation

$$\begin{aligned}P(x, C_1) &= P(x|C_1)P(C_1) \\&= \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1) \right\} P(C_1)\end{aligned}$$

$$\begin{aligned}P(x, C_2) &= P(x|C_2)P(C_2) \\&= \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2}(x - \mu_2)^T \Sigma^{-1}(x - \mu_2) \right\} P(C_2)\end{aligned}$$

$$\log \frac{P(C_1|x)}{P(C_2|x)} = \log \frac{P(C_1|x)}{1 - P(C_1|x)} \quad \text{“Log-odds”}$$

$$\begin{aligned}&= \log \frac{\exp \left\{ -\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1) \right\}}{\exp \left\{ -\frac{1}{2}(x - \mu_2)^T \Sigma^{-1}(x - \mu_2) \right\}} + \log \frac{P(C_1)}{P(C_2)} \\&= \left\{ -\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1) \right\} - \left\{ -\frac{1}{2}(x - \mu_2)^T \Sigma^{-1}(x - \mu_2) \right\} + \log \frac{P(C_1)}{P(C_2)} \\&= (\mu_1 - \mu_2)^T \Sigma^{-1}x - \frac{1}{2}\mu_1^T \Sigma^{-1}\mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1}\mu_2 + \log \frac{P(C_1)}{P(C_2)} \\&= (\Sigma^{-1}(\mu_1 - \mu_2))^T x + w_0\end{aligned}$$

Quadratic term canceled out because of the shared covariance

$$\text{where } w_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1}\mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1}\mu_2 + \log \frac{P(C_1)}{P(C_2)}$$


# Extension to Multi-Class GDA

- For two-class classification, the posterior  $p(C_k|\mathbf{x})$  is the sigmoid of the log odds

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

$$a = \log \left( \frac{\sigma}{1-\sigma} \right) = (\Sigma^{-1}(\mu_1 - \mu_2))^T x + w_0$$

$$\text{where } w_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \log \frac{P(C_1)}{P(C_2)}$$





- For multi-class classification, the posterior  $p(C_k|\mathbf{x})$  is softmax.

$$p_i = \frac{\exp(q_i)}{\sum_j \exp(q_j)}$$

# GDA Decision Boundaries

- At decision boundary, we have  $p(C_1|\mathbf{x}) = p(C_2|\mathbf{x})$
- Under the same covariance assumption, the boundary is linear.
  - Different priors  $p(C_1), p(C_2)$  does not change the linearity but shift it around.





# Learning GDA via MLE

- Given training data  $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$  and a generative model with the shared covariance
- Priors:  $p(y) = \phi^y (1 - \phi)^{1-y}$  where  $\phi = p(y = 1)$
- Class densities:

$$p(\mathbf{x}|y = 0) = \frac{1}{\sqrt{2\pi} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu_0)^T \Sigma^{-1}(\mathbf{x} - \mu_0)\right)$$

$$p(\mathbf{x}|y = 1) = \frac{1}{\sqrt{2\pi} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu_1)^T \Sigma^{-1}(\mathbf{x} - \mu_1)\right)$$

# Learning GDA via MLE

- Maximum likelihood estimation (MLE):

$$\phi = \frac{1}{N} \sum_{i=1}^N \mathbf{1}\{y^{(i)} = 1\}$$



$$\mu_0 = \frac{\sum_{i=1}^N \mathbf{1}\{y^{(i)} = 0\} \mathbf{x}^{(i)}}{\sum_{i=1}^N \mathbf{1}\{y^{(i)} = 0\}}$$

$$\mu_1 = \frac{\sum_{i=1}^N \mathbf{1}\{y^{(i)} = 1\} \mathbf{x}^{(i)}}{\sum_{i=1}^N \mathbf{1}\{y^{(i)} = 1\}}$$

$$\sum = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \mu_{y^{(i)}})(\mathbf{x}^{(i)} - \mu_{y^{(i)}})^T$$

# GDA with Different Covariance

- Decision boundaries are quadratic when each class has different covariance.



# GDA vs. Logistic Regression

- GDA requires  $O(M^2)$  learnable parameters
  - $2M$  parameters for the means of  $p(\mathbf{x}|C_1)$  and  $p(\mathbf{x}|C_2)$
  - $M(M + 1)/2$  parameters for shared covariance matrix
  - Cf. Logistic regression requires only  $M$  parameters
- GDA has a strong modeling assumption and works well when the distribution follows the assumption.
  - Cf. Logistic regression has less parameters and is more flexible about data distribution.

# Naïve Bayes Classifier

# Naïve Bayes Classifier

- Prior distribution  $p(C_k)$ : Constant (e.g., Bernoulli)
- Likelihood  $p(\mathbf{x}|C_k)$ :

$$P(x_1, \dots, x_M | C_k) = P(x_1 | C_k) \cdots P(x_M | C_k) = \prod_{j=1}^M P(x_j | C_k)$$

- Naïve Bayes assumption: Each coordinate of  $\mathbf{x}$  is conditionally independent of other coordinates given the class label.
- Classification using Bayes' rule:

$$p(C_k | \mathbf{x}) = p(\mathbf{x} | C_k) p(C_k) / p(\mathbf{x})$$

- For two-class classification,

$$P(C_1 | \mathbf{x}) = \frac{P(C_1, \mathbf{x})}{P(\mathbf{x})} = \frac{P(C_1, \mathbf{x})}{P(C_1, \mathbf{x}) + P(C_2, \mathbf{x})}$$

# Naïve Bayes Classifier

- Classification using Bayes' rule:

$$p(C_k|\mathbf{x}) = p(\mathbf{x}|C_k)p(C_k)/p(\mathbf{x})$$

- Classification is done by the MAP estimation:

$$\arg \max_k P(C_k|\mathbf{x}) = \arg \max_k P(C_k, \mathbf{x})$$

Naïve Bayes  
Assumption




$$= \arg \max_k P(C_k)P(\mathbf{x}|C_k)$$

$$= \arg \max_k P(C_k) \prod_{j=1}^M P(x_j|C_k)$$

# Example: Spam Mail Classification

- Label:  $y = 1$  (spam),  $y = 0$  (ham or non-spam)
- Features  $\mathbf{x} = [x_1, x_2, \dots]$ 
  - $x_k$ :  $k$ -th word in a mail, where  $M$  is the vocabulary size
  - Each word is represented as **one-hot encoding**.



- Naïve Bayes assumption: Given a class label  $y$ , each word in a mail is an independent **multinomial** variable.



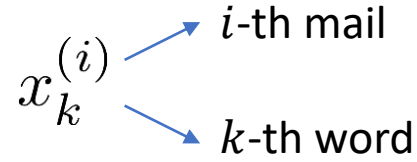
# Naïve Bayes Classifier Formulation

- Prior:  $P(\text{spam}) = \text{Bernoulli}(\phi)$
- Likelihood:  $P(\text{word}|\text{spam}) = \text{Multinomial}(\mu_1^s, \dots, \mu_M^s)$   
 $P(\text{word}|\text{nospam}) = \text{Multinomial}(\mu_1^{ns}, \dots, \mu_M^{ns})$
- Learning to find  $\phi, \mu^s, \mu^{ns}$  that best fits the training data  $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$

$$\begin{aligned} & \prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)}) \\ &= \prod_{i=1}^N P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \\ &= \left( \prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \left( \prod_{i:y^{(i)}=0} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \\ & \quad \text{Spam} \qquad \qquad \text{Ham (Non-spam)} \end{aligned}$$

# Naïve Bayes Classifier Derivation

- Likelihood for spam:

$$\left( \prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right)$$


$x_k^{(i)}$   $\rightarrow$   $i$ -th mail  
 $x_k^{(i)}$   $\rightarrow$   $k$ -th word

- Naïve Bayes assumption:

- Prior:  $P(\text{spam}) = \text{Bernoulli}(\phi)$

$$P(y^{(i)} = 1) = \phi$$

- Likelihood:  $P(\text{word}|\text{spam}) = \text{Multinomial}(\mu_1^s, \dots, \mu_M^s)$

$$\begin{aligned} P(x^{(i)}|y^{(i)} = 1) &= \prod_{k=1}^{\text{len}(x^{(i)})} P(x_k^{(i)}|y^{(i)} = 1) \\ &= \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \end{aligned}$$

# Naïve Bayes Classifier Derivation


- Likelihood for spam:

$$\begin{aligned} & \left( \prod_{i: y^{(i)}=1} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \\ &= \left( \prod_{i: y^{(i)}=1} \left( \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = j\text{-th word})} \right) \phi \right) \end{aligned}$$

# Naïve Bayes Classifier Derivation

- Likelihood for spam:


$$\begin{aligned}
 & \left( \prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \\
 = & \left( \prod_{i:y^{(i)}=1} \left( \prod_{k=1}^{len(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \right) \phi \right) \\
 = & \left( \prod_{i:y^{(i)}=1} \prod_{k=1}^{len(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \right) \left( \prod_{i:y^{(i)}=1} \phi \right)
 \end{aligned}$$


 $\prod_i a_i b = \prod_i a_i \prod_i b$

# Naïve Bayes Classifier Derivation

- Likelihood for spam:

$$\begin{aligned}
 & \left( \prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \\
 &= \left( \prod_{i:y^{(i)}=1} \left( \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \right) \phi \right) \\
 &= \left( \prod_{i:y^{(i)}=1} \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \right) \left( \prod_{i:y^{(i)}=1} \phi \right) \\
 &= \left( \prod_{j=1}^M (\mu_j^s)^{\sum_{i:y^{(i)}=1} \sum_{k=1}^{\text{len}(x^{(i)})} I(x_k^{(i)}=j\text{-th word})} \right) \left( \prod_{i:y^{(i)}=1} \phi \right)
 \end{aligned}$$


 $\prod_i a_j^b = a_j^{\sum_i b}$

# Naïve Bayes Classifier Derivation

- Likelihood for spam:

$$\begin{aligned}
 & \left( \prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \\
 = & \left( \prod_{i:y^{(i)}=1} \left( \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \right) \phi \right) \\
 = & \left( \prod_{i:y^{(i)}=1} \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \right) \left( \prod_{i:y^{(i)}=1} \phi \right) \\
 = & \left( \prod_{j=1}^M (\mu_j^s)^{\sum_{i:y^{(i)}=1} \sum_{k=1}^{\text{len}(x^{(i)})} I(x_k^{(i)}=j\text{-th word})} \right) \left( \prod_{i:y^{(i)}=1} \phi \right) \\
 = & \left( \prod_{j=1}^M (\mu_j^s)^{N_j^{\text{spam}}} \right) \phi^{N^{\text{spam}}}
 \end{aligned}$$

For  $i$ -th email,

count # of  $j$ -th word in the vocabulary


$N_j^{\text{spam}}$ : Total # of  $j$ -th word in spam emails

$N^{\text{spam}}$ : Total # of spam emails

# Naïve Bayes Classifier Derivation

- Likelihood for spam:

$$\left( \prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right)$$



- Naïve Bayes assumption:

- Prior:  $P(\text{spam}) = \text{Bernoulli}(\phi)$

$$P(y^{(i)} = 1) = \phi$$

- Likelihood:  $P(\text{word}|\text{spam}) = \text{Multinomial}(\mu_1^s, \dots, \mu_M^s)$

$$\begin{aligned} P(x^{(i)}|y^{(i)} = 1) &= \prod_{k=1}^{\text{len}(x^{(i)})} P(x_k^{(i)}|y^{(i)} = 1) \\ &= \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)}=j\text{-th word})} \end{aligned}$$

# Naïve Bayes Classifier Derivation

- Likelihood for ham (non-spam):

$$\left( \prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \left( \prod_{i:y^{(i)}=0} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right)$$

- Naïve Bayes assumption:

- Prior:  $P(\text{spam}) = \text{Bernoulli}(\phi)$

$$P(y^{(i)} = 0) = 1 - \phi$$

- Likelihood:  $P(\text{word}|\text{nospam}) = \text{Multinomial}(\mu_1^{ns}, \dots, \mu_M^{ns})$

$$\begin{aligned} P(x^{(i)}|y^{(i)} = 0) &= \prod_{k=1}^{\text{len}(x^{(i)})} P(x_k^{(i)}|y^{(i)} = 0) \\ &= \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^{ns})^{I(x_k^{(i)} = j\text{-th word})} \end{aligned}$$



# Naïve Bayes Classifier Derivation

- Putting together:

$$\begin{aligned} & \prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)}) \\ &= \left( \prod_{i: y^{(i)}=1} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \left( \prod_{i: y^{(i)}=0} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \\ &= \left( \phi^{N^{spam}} \prod_{word\ j} (\mu_j^s)^{N_j^{spam}} \right) \left( (1 - \phi)^{N^{nonspam}} \prod_{word\ j} (\mu_j^{ns})^{N_j^{nonspam}} \right) \end{aligned}$$

- Log-likelihood

$$\begin{aligned} & \log P(\mathcal{D}) \\ &= \log \prod_{i=1}^N P(x^{(i)}, y^{(i)}) \\ &= N^{spam} \log \phi + \sum_{word\ j} N_j^{spam} \log \mu_j^s + N^{nonspam} \log(1 - \phi) + \sum_{word\ j} N_j^{nonspam} \log \mu_j^{ns} \end{aligned}$$

# Learning NB Classifier via MLE

- Log-likelihood

$$\begin{aligned} & \log P(\mathcal{D}) \\ = & \log \prod_{i=1}^N P(x^{(i)}, y^{(i)}) \\ = & N^{spam} \log \phi + \sum_{word\ j} N_j^{spam} \log \mu_j^s + N^{nonspam} \log(1 - \phi) + \sum_{word\ j} N_j^{nonspam} \log \mu_j^{ns} \end{aligned}$$

- Maximum (log-)likelihood estimation

- Take the derivative of log-likelihood with respect to the parameters  $\{\phi, \mu^s, \mu^{ns}\}$ , and set it to zero.

# Learning NB Classifier via MLE

- Find  $\phi$

$$\log P(\mathcal{D})$$

$$= N^{spam} \log \phi + \sum_{word\ j} N_j^{spam} \log \mu_j^s + N^{nonsпам} \log(1 - \phi) + \sum_{word\ j} N_j^{nonsпам} \log \mu_j^{ns}$$

$$\Rightarrow \frac{\partial l}{\partial \phi} = \frac{1}{\phi} N^{spam} - \frac{1}{1 - \phi} N^{nonsпам} = 0$$

$$\phi = \frac{N^{spam}}{N^{spam} + N^{nonsпам}}$$

# Learning NB Classifier via MLE

- Find  $\mu^s$  (or similarly,  $\mu^{ns}$ )
  - $\{\mu^s\}$ 's are NOT independent to each other;  $\sum_{j=1}^M \mu_j^s = 1$
  - To only deal with variables independent to each other, we can express  $\mu_M^s$  as  $1 - \sum_{j=1}^{M-1} \mu_j^s$



$$\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s = \sum_{word\ j=1}^{M-1} N_j^{spam} \log \mu_j^s + N_M^{spam} \log(1 - \sum_{j=1}^{M-1} \mu_j^s)$$

# Learning NB Classifier via MLE

- Find  $\mu^s$  (or similarly,  $\mu^{ns}$ )
  - $\{\mu^s\}$ 's are NOT independent to each other;  $\sum_{j=1}^M \mu_j^s = 1$
  - To only deal with variables independent to each other, we can express  $\mu_M^s$  as  $1 - \sum_{j=1}^{M-1} \mu_j^s$

$$\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s = \sum_{word\ j=1}^{M-1} N_j^{spam} \log \mu_j^s + N_M^{spam} \log(1 - \sum_{j=1}^{M-1} \mu_j^s)$$

Take derivative

$$\frac{\partial}{\partial \mu_j^s} \left( \sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s \right) = \frac{N_j^{spam}}{\mu_j^s} - \boxed{\frac{N_M^{spam}}{1 - \sum_{j=1}^{M-1} \mu_j^s}} = 0$$

Constant w.r.t.  $j$

# Learning NB Classifier via MLE

- Find  $\mu^s$  (or similarly,  $\mu^{ns}$ )
  - $\{\mu^s\}$ 's are NOT independent to each other;  $\sum_{j=1}^M \mu_j^s = 1$
  - To only deal with variables independent to each other, we can express  $\mu_M^s$  as  $1 - \sum_{j=1}^{M-1} \mu_j^s$

$$\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s = \sum_{word\ j=1}^{M-1} N_j^{spam} \log \mu_j^s + N_M^{spam} \log(1 - \sum_{j=1}^{M-1} \mu_j^s)$$

$$\frac{\partial}{\partial \mu_j^s} \left( \sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s \right) = \frac{N_j^{spam}}{\mu_j^s} - \boxed{\frac{N_M^{spam}}{1 - \sum_{j=1}^{M-1} \mu_j^s}} = 0$$

Constant w.r.t.  $j$

$$\frac{a}{b} = \frac{c}{d} = \frac{a+c}{b+d}$$

$$\frac{N_j^{spam}}{\mu_j^s} = \text{constant} = \frac{\sum_{j=1}^M N_j^{spam}}{\sum_{j=1}^M \mu_j^s} = \sum_{j=1}^M N_j^{spam} = N^{spam}$$

# Learning NB Classifier via MLE

- Find  $\mu^s$  (or similarly,  $\mu^{ns}$ )
  - $\{\mu^s\}$ 's are NOT independent to each other;  $\sum_{j=1}^M \mu_j^s = 1$
  - To only deal with variables independent to each other, we can express  $\mu_M^s$  as  $1 - \sum_{j=1}^{M-1} \mu_j^s$

$$\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s = \sum_{word\ j=1}^{M-1} N_j^{spam} \log \mu_j^s + N_M^{spam} \log(1 - \sum_{j=1}^{M-1} \mu_j^s)$$

$$\frac{\partial}{\partial \mu_j^s} \left( \sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s \right) = \frac{N_j^{spam}}{\mu_j^s} - \frac{N_M^{spam}}{1 - \sum_{j=1}^{M-1} \mu_j^s} = 0$$

$$\therefore \mu_j^s = \frac{N_j^{spam}}{\sum_j N_j^{spam}}$$

# Learning NB Classifier via MLE

- Summary

$$P(spam) = \phi = \frac{N^{spam}}{N^{spam} + N^{nonspam}}$$

$$P(word = j|spam) = \mu_j^s = \frac{N_j^{spam}}{\sum_j N_j^{spam}}$$

$$P(word = j|non - spam) = \mu_j^{ns} = \frac{N_j^{nonspam}}{\sum_j N_j^{nonspam}}$$

- $N^{spam}$ : Total # of spam emails
- $N^{nonspam}$ : Total # of ham (non-spam) emails
- $N_j^{spam}$ : Total # of  $j$ -th word in spam emails
- $N_j^{nonspam}$ : Total #  $j$ -th word in ham (non-spam) emails



# Laplace Smoothing

- Maximum likelihood is problematic when a specific word count is 0.
  - Leads to probability of the specific word 0
- Solution: Put **imaginary** counts for each word
  - Prevent zero probability estimates (overfitting)
  - Add “1” as imaginary count for each word

$$P(spam) = \phi = \frac{N^{spam}}{N^{spam} + N^{nonspam}}$$

$$P(word = j|spam) = \mu_j^s = \frac{N_j^{spam} + 1}{\sum_j N_j^{spam} + M}$$

$$P(word = j|non - spam) = \mu_j^{ns} = \frac{N_j^{nonspam} + 1}{\sum_j N_j^{nonspam} + M}$$

# Laplace Smoothing

- Maximum likelihood is problematic when a specific word count is 0.
  - Leads to probability of the specific word 0
- Solution: Put **imaginary** counts for each word
  - Prevent zero probability estimates (overfitting)
  - Add “1” as imaginary count for each word

$$P(spam) = \phi = \frac{N^{spam}}{N^{spam} + N^{nonspam}}$$

- If we smooth prior as well:

$$P(spam) = \phi = \frac{N^{spam} + 1}{N^{spam} + N^{nonspam} + 2}$$

- (Don't have to do this)

# Next: Other Classifiers