

12. Kernel Methods

STA3142 Statistical Machine Learning

Kibok Lee

Assistant Professor of

Applied Statistics / Statistics and Data Science

Apr 9, 2024



연세대학교
YONSEI UNIVERSITY

Assignment 2

- Due **Friday 4/12, 11:59pm**
- Topics
 - (Math/Programming) Logistic Regression
 - (Math/Programming) Softmax Regression
 - (Math) Gaussian Discriminant Analysis
 - (Programming) Naïve Bayes for Spam Classification
- Please read the instruction carefully!
 - Submit one pdf and one zip file separately
 - Write your code only in the designated spaces
 - Do not import additional libraries
 - ...
- If you feel difficult, consider to take **option 2**.

Midterm

- **Tuesday 4/23, 1:10pm — 2:50pm KST**
 - Please come here by 1:00pm!
 - In-person exam
- Closed book with **an A4-size cheat sheet**
 - You can print/write anything on **both side**.
- Coverage: Lec 6—13
 - True / False, multiple choice, math
- Short practice midterm will be out.
 - To be familiar with the type of midterm questions
 - # questions is about a half of the actual exam
 - **No solution will be provided**

Midterm Coverage


- 4,5: Linear Algebra & Probability Review
 - Not main topics, but you should be familiar with them.
 - Some contents (that we feel difficult) can be given FYI.
- 6,7. Linear Regression (and Other Topics)
- 8. Logistic/Softmax Regression
- 9. Generative Classifiers
- 10. Other Classifiers
- 11. Regularization and Validation
- 12. Kernel Methods
- 13. Support Vector Machines

Announcement

- 24 Fall graduate school application
 - Submission: 4/19 ~ 4/26, 5 PM
 - 1st notification: 5/24, 5 PM
 - Exam: 6/1
 - Final notification: 6/14, 5 PM
 - Registration: July ~ Aug


**2024학년도 후기
일반대학원 신입생 모집 안내**

원서 및 서류 제출(온라인)	2024. 4. 19.(금) ~ 4. 26.(금) 17:00 마감
구술/실기시험 대상자 발표	2024. 5. 24.(금) 17:00
구술/실기시험	2024. 6. 1.(토)
최종합격자 발표	2024. 6. 14.(금) 17:00
등록	2024년 7~8월 예정 (2024년 9월 신입학)

 연세대학교 대학원

일반대학원 홈페이지(graduate.yonsei.ac.kr)
"입학안내 및 팝업" 참고

연세대 일반대학원




- More Information/Submission:
 - Kor: <https://graduate.yonsei.ac.kr/graduate/index.do>
 - Eng: https://graduate.yonsei.ac.kr/graduate_en/index.do


Recap: Regularization

- Regularization controls the tradeoff between “fitting error” and “complexity.”
 - Small regularization results in complex models (with risk of overfitting)
 - Large regularization results in simple models (with risk of underfitting)


Overfitting



Sweet spot





Underfitting



Recap: The Bias-Variance Tradeoff

- An over-regularized model (large λ) will have a **high bias** and **low variance**.
- An under-regularized model (small λ) will have a **high variance** and **low bias**.
- It is important to find a good balance between the two.



Recap: Validation

- If model selection and true error estimates are to be computed simultaneously, data needs to be divided into three disjoint sets.
- **Training set** to fit the parameters
 - Given a fixed hyperparameters
- **Validation set** to tune/choose the model and hyperparameters
- **Test set** to evaluate the final model performance
 - You must **NOT** tune the model on test set.
 - Test set is **NOT** for model selection.



Recap: K-Fold Cross-Validation

- Split dataset into K-folds
 - Take one fold (yellow) as validation and the rest of K-1 folds (green) for training.

Trial 1	fold 1	fold 2	fold 3	fold 4
Trial 2	fold 1	fold 2	fold 3	fold 4
Trial 3	fold 1	fold 2	fold 3	fold 4
Trial 4	fold 1	fold 2	fold 3	fold 4


- The final validation error is estimated as the average error rate.

Outline

- Feature Mappings
- Kernel Trick
- Dual Representations
 - Example: Kernel Ridge Regression
- Constructing Kernels
- (Nadaraya-Watson) Kernel Regression

Linear Regression

- Linear function $h(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x}$ can only produce straight lines through origin.
- Affine function $h(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x} + b$ can only produce straight lines.
 - e.g., 1D linear regression
 - Not very flexible/powerful




- Can we make it more flexible?


Feature Mappings

- Solution: Add features and hope some of them are useful; e.g., polynomial basis functions

- $\phi(x) = (1, x)$



- $\phi(x) = (1, x, x^2, x^3)$



Linear Regression with Features

- Linear regression model

$$h(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \phi(\mathbf{x}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x})$$


- Least squares with L2 regularization

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^T \phi(\mathbf{x}_n) - y_n)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- Gradient: $\nabla_{\mathbf{w}} J(\mathbf{w}) = \Phi^T (\Phi \mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}$
- Closed form solution: $\mathbf{w} = (\Phi^T \Phi + \lambda I)^{-1} \Phi^T \mathbf{y}$

Linear Classifiers


- Linear function $h(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x}$ can only produce linear decision boundaries through origin.
- Affine function $h(\mathbf{x}; \mathbf{w}) = \mathbf{w}^T \mathbf{x} + b$ can only produce linear decision boundaries.
 - e.g., 2D linear classification
 - Not very flexible/powerful




- Can we make it more flexible?

Linear Classifiers with Features

- $\phi(x_1, x_2) = (x_1^2, x_2^2)$



Not linearly separable



Linearly separable

- (Nonlinear) features can make the problem solvable with linear methods.

Feature Selection

- We have been mapping data \mathbf{x} through a fixed (nonlinear) mapping to get a feature vector $\phi(\mathbf{x})$.
 - The feature vector extracts important properties from \mathbf{x} .
 - e.g., Polynomial combinations up to some order
 - It makes many regression/classification problems easier.
- In general, what features should we use?

Feature Selection

- How about to use **all features** we can think of?
- Model complexity might not be an issue, as proper regularization can handle it.
- However, it is **not scalable**; assuming large $N \approx M$,
 - N : Number of training data
 - M : Number of features
- Closed-form solution requires $O(N^3)$
- Gradient descent requires $O(N^2)$

Feature Selection

- We have been mapping data \mathbf{x} through a fixed (nonlinear) mapping to get a feature vector $\phi(\mathbf{x})$.
- In other words, with feature mappings, all data have been mapped to a higher dimensional space.
- Alternatively, we can think of it like:
Data still lives in the original space, but the definition of **distance** or **inner product** has been changed.

Kernel Trick

Kernel Trick

- As we have done, we will embed data \mathbf{x} in a high dimensional space $\phi(\mathbf{x})$, and use simple (linear) models in this space.
- Use algorithms that do not need the coordinates of embeddings $\phi(\mathbf{x})$, but pairwise **inner products**:

$$\phi(\mathbf{x})^T \phi(\mathbf{x}')$$

- Replace these inner products with a **kernel**:

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$

Kernel Trick

- A **kernel** function $k(\mathbf{x}, \mathbf{x}')$ represents the similarity between \mathbf{x} and \mathbf{x}' .

- A popular way to express the similarity between feature vectors is the **inner product** of them:

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$$

- A **kernel** function $k(\mathbf{x}, \mathbf{x}')$ is defined to be an inner product of feature vectors, but **we do not actually compute them.**

Example: Kernels for 2D Data

- Inner product between (x_1, x_2) and (z_1, z_2) :

$$k(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z} = x_1 z_1 + x_2 z_2$$

- Its square is also a kernel:

$$\begin{aligned} k(\mathbf{x}, \mathbf{z}) &= (\mathbf{x}^T \mathbf{z})^2 = (x_1 z_1 + x_2 z_2)^2 \\ &= x_1^2 z_1^2 + 2x_1 z_1 x_2 z_2 + x_2^2 z_2^2 \end{aligned}$$

- This is the same as inner product between

$$(x_1^2, \sqrt{2}x_1 x_2, x_2^2) \text{ and } (z_1^2, \sqrt{2}z_1 z_2, z_2^2)$$

- Or, between $(x_1^2, x_1 x_2, x_1 x_2, x_2^2)$ and $(z_1^2, z_1 z_2, z_1 z_2, z_2^2)$

- **Note: Solution is not unique.**

Example: Kernels for 2D Data

- Consider higher-order polynomial of degree p :

$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z})^p = \left(\sum_{i=1}^M x_i z_i \right)^p$$

← Multinomial expansion

$$= \sum_{(j_1, \dots, j_M): \sum_j j_k = p} \binom{p}{j_1 \ j_2 \ \dots \ j_M} (x_1 z_1)^{j_1} \dots (x_M z_M)^{j_M}$$

- Feature mapping:

$$\phi(\mathbf{x}) = \left[\dots, \binom{p}{j_1 \ j_2 \ \dots \ j_M}^{\frac{1}{2}} (x_1)^{j_1} \dots (x_M)^{j_M}, \dots \right]$$

- All monomials of degree p

Example: Kernels for 2D Data

- Inhomogeneous polynomial up to degree p :


$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z} + c)^p = \left(c + \sum_{i=1}^M x_i z_i \right)^p, c > 0$$

- Feature mapping:

$\phi(\mathbf{x}) =$ all monomials of degree $\leq p$.

Example: Handwritten Digits

- An image consists of $28 \times 28 = 784$ pixels
 $\mathbf{x} \in [0,1]^{784}$



- Take the pixel values and compute
$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z} + 1)^p$$
- For $p = 4$, computing the inner product in the space of all monomials **without kernel trick** requires **16G** dimensional space.

Kernel Trick

- With kernels, inner products in a high dimensional space can be computed with the computational complexity of a low dimensional space.
- Many algorithms can be expressed completely in terms of kernels $k(\mathbf{x}, \mathbf{x}')$ without $\phi(\mathbf{x})$.
- We can replace a kernel with another to get a new algorithm that works over a different domain.

Kernel Trick

- To use the kernel trick, we must formulate (training and test) algorithms purely in terms of inner products between data points.
 - All operations w.r.t. \mathbf{x} should look like $k(\mathbf{x}, \mathbf{x}')$
- We cannot access the coordinates in the high-dimensional feature space, i.e., no explicit $\phi(\mathbf{x})$.
- This seems a huge limitation, but many operations/algorithms work under this condition.

Example: Distance

- Distance between two samples can be expressed in inner products:

$$\begin{aligned} & \|\phi(\mathbf{x}) - \phi(\mathbf{z})\|^2 \\ &= \phi(\mathbf{x})^T \phi(\mathbf{x}) - 2\phi(\mathbf{x})^T \phi(\mathbf{z}) + \phi(\mathbf{z})^T \phi(\mathbf{z}) \\ &= k(\mathbf{x}, \mathbf{x}) - 2k(\mathbf{x}, \mathbf{z}) + k(\mathbf{z}, \mathbf{z}) \end{aligned}$$

- e.g., K-nearest neighbors can be done in an arbitrary high dimensional space with the kernel trick.

Example: Mean

- Mean of data points:

$$\bar{\phi} = \frac{1}{N} \sum_{i=1}^N \phi(\mathbf{x}^{(i)})$$

- We cannot determine the mean of data in the mapped feature space through kernel operations.

Example: Distance to Mean

- Mean of data points:

$$\bar{\phi} = \frac{1}{N} \sum_{i=1}^N \phi(\mathbf{x}^{(i)})$$

- Distance to mean:

$$\begin{aligned} & \|\phi(\mathbf{x}) - \bar{\phi}\|^2 \\ &= \phi(\mathbf{x})^T \phi(\mathbf{x}) + \bar{\phi}^T \bar{\phi} - 2\phi(\mathbf{x})^T \bar{\phi} \\ &= k(\mathbf{x}, \mathbf{x}) + \frac{1}{N^2} \sum_{i,j=1}^N k(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) - \frac{2}{N} \sum_{i=1}^N k(\mathbf{x}, \mathbf{x}^{(i)}) \end{aligned}$$

Dual Representations

Dual Representations and Kernel Trick

- The dual representation and its solution are entirely written in terms of kernels.

- The elements of the Gram matrix $K = \Phi\Phi^T$

$$K_{ij} = \phi(\mathbf{x}^{(i)})^T \phi(\mathbf{x}^{(j)}) = k(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$$

- These represent the pairwise similarities among all the observed feature vectors.
 - We can compute kernels more efficiently than the feature vectors.

Example: Kernel Ridge Regression

- Recall regression problems with error function

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^T \phi(\mathbf{x}^{(n)}) - y^{(n)})^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- $J(\mathbf{w})$ is minimized at

$$\mathbf{w} = (\Phi^T \Phi + \lambda I)^{-1} \Phi^T \mathbf{y}$$

- Recall the $N \times M$ data matrix that is central to this solution.

Recap: The Data Matrix

- The data matrix is an $N \times M$ matrix, applying
 - the M basis functions (columns)
 - to N data points (rows)

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}^{(1)}) & \phi_1(\mathbf{x}^{(1)}) & \dots & \phi_{M-1}(\mathbf{x}^{(1)}) \\ \phi_0(\mathbf{x}^{(2)}) & \phi_1(\mathbf{x}^{(2)}) & \dots & \phi_{M-1}(\mathbf{x}^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}^{(N)}) & \phi_1(\mathbf{x}^{(N)}) & \dots & \phi_{M-1}(\mathbf{x}^{(N)}) \end{pmatrix}$$

$$\Phi \mathbf{w} \approx \mathbf{y}$$

The Gram Matrix

- For closed-form solution, we use the $M \times M$ (scaled covariance) matrix: $\Phi^T \Phi$
- For dual representation, we use the $N \times N$ Gram matrix: $K = \Phi \Phi^T$
 - where $K_{ij} = \phi(\mathbf{x}^{(i)})^T \phi(\mathbf{x}^{(j)}) = k(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$, each element corresponds to the pairwise similarities of two training data.
- Note that kernel methods use only K , not Φ .

Example: Kernel Ridge Regression

- Objective function:

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^T \phi(\mathbf{x}^{(n)}) - y^{(n)})^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- Take derivative: $\nabla_{\mathbf{w}} J(\mathbf{w}) = \Phi^T (\Phi \mathbf{w} - \mathbf{y}) + \lambda \mathbf{w} = 0$
- Then,

$$\mathbf{w} = -\frac{1}{\lambda} \Phi^T (\Phi \mathbf{w} - \mathbf{y}) = \Phi^T \mathbf{a} = \sum_{n=1}^N a_n \phi(\mathbf{x}^{(n)})$$

- where $\mathbf{a} = [a_1, \dots, a_N]^T$; $a_n = -\frac{1}{\lambda} (\mathbf{w}^T \phi(\mathbf{x}^{(n)}) - y^{(n)})$
be the new parameters.
- i.e., transform $J(\mathbf{w})$ to $J(\mathbf{a})$ with $\mathbf{w} = \Phi^T \mathbf{a}$

Example: Kernel Ridge Regression

- Objective function:

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (\mathbf{w}^T \phi(\mathbf{x}^{(n)}) - y^{(n)})^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- Substitute $\mathbf{w} = \Phi^T \mathbf{a}$:

$$\begin{aligned} J(\mathbf{a}) &= \frac{1}{2} \mathbf{a}^T \Phi \Phi^T \Phi \Phi^T \mathbf{a} - \mathbf{a}^T \Phi \Phi^T \mathbf{y} + \frac{1}{2} \mathbf{y}^T \mathbf{y} + \frac{\lambda}{2} \mathbf{a}^T \Phi \Phi^T \mathbf{a} \\ &= \frac{1}{2} \mathbf{a}^T K K \mathbf{a} - \mathbf{a}^T K \mathbf{y} + \frac{1}{2} \mathbf{y}^T \mathbf{y} + \frac{\lambda}{2} \mathbf{a}^T K \mathbf{a} \end{aligned}$$

- Solution: $\mathbf{a} = (K + \lambda I_N)^{-1} \mathbf{y}$

- At test time:

$$h(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) = \mathbf{a}^T \Phi \phi(\mathbf{x}) = \mathbf{k}(\mathbf{x})^T (K + \lambda I_N)^{-1} \mathbf{y}$$

- where $\mathbf{k}(\mathbf{x}) = [k(\mathbf{x}^{(1)}, \mathbf{x}), \dots, k(\mathbf{x}^{(N)}, \mathbf{x})]^T$

Primal vs. Dual Representations

- Primal: $\mathbf{w} = (\Phi^T \Phi + \lambda I_M)^{-1} \Phi^T \mathbf{y}$
 - Need to invert an $M \times M$ matrix, where M is the feature dimension
 - Efficient when $N > M$
 - Need to compute features explicitly
- Dual: $\mathbf{a} = (K + \lambda I_N)^{-1} \mathbf{y}$
 - Need to invert an $N \times N$ matrix, where N is the number of training data
 - Efficient when $N < M$
 - Kernel trick is applicable; kernels can be defined over vectors, images, sequences, graphs, text, etc.

Constructing Kernels

Constructing Kernels: Method 1

- We can do kernel engineering to create kernels for particular purposes, expressing different kinds of similarity.
- Define a feature mapping $\phi(\mathbf{x})$ and then define the inner product of features as kernel (or vice versa)
 - Formally, $k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is a valid kernel if and only if there exists $\phi: \mathcal{X} \rightarrow \mathcal{H}$ such that
$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}'), \forall \mathbf{x}, \mathbf{x}' \in \mathcal{X}$$
 - where \mathcal{H} is a Hilbert space (= Euclidean space with potentially infinite-dimensional)

Constructing Kernels: Method 1

- Define a kernel function directly, such as

$$k(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z} = x_1 z_1 + x_2 z_2$$

- In 2D, we can explicitly identify the feature map

$$\phi(\mathbf{x}) = (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$

- such that $k(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x})^T \phi(\mathbf{z})$
- Explicit feature mappings can be very complex.
 - Kernels help us avoid that complexity.

Constructing Kernels: Method 2

- A simpler way to test if a given function is kernel without constructing $\phi(\mathbf{x})$
1. Show that the Gram matrix K is positive semidefinite (PSD) for all possible choices of the dataset.

$$\mathbf{a}^T K \mathbf{a} \equiv \sum_{ij} a_i K_{ij} a_j \geq 0, \forall \mathbf{a} \in R^N$$

2. Use **Mercer's theorem** to prove that the kernel function $k(\mathbf{x}, \mathbf{x}')$ is a valid kernel
 - Mercer's theorem is **the necessary and sufficient condition**: K is PSD iff $k(\mathbf{x}, \mathbf{x}')$ is a valid kernel

Constructing Kernels: Method 3

- There are a number of axioms that help us construct new, more complex kernels, from simpler known kernels.
- For example, with a known kernel k_1 ,

$$k(\mathbf{x}, \mathbf{x}') = f(\mathbf{x})k_1(\mathbf{x}, \mathbf{x}')f(\mathbf{x}')$$

$$k(\mathbf{x}, \mathbf{x}') = \exp(k_1(\mathbf{x}, \mathbf{x}'))$$

$$k(\mathbf{x}, \mathbf{x}') = \exp(-||\mathbf{x} - \mathbf{x}'||^2/2\sigma^2)$$

Constructing Kernels: Method 3

Given valid kernels $k_1(\mathbf{x}, \mathbf{x}')$ and $k_2(\mathbf{x}, \mathbf{x}')$, the following new kernels will also be valid:

$$k(\mathbf{x}, \mathbf{x}') = ck_1(\mathbf{x}, \mathbf{x}') \quad (6.13)$$

$$k(\mathbf{x}, \mathbf{x}') = f(\mathbf{x})k_1(\mathbf{x}, \mathbf{x}')f(\mathbf{x}') \quad (6.14)$$

$$k(\mathbf{x}, \mathbf{x}') = q(k_1(\mathbf{x}, \mathbf{x}')) \quad (6.15)$$

$$k(\mathbf{x}, \mathbf{x}') = \exp(k_1(\mathbf{x}, \mathbf{x}')) \quad (6.16)$$

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}') \quad (6.17)$$

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}') \quad (6.18)$$

$$k(\mathbf{x}, \mathbf{x}') = k_3(\phi(\mathbf{x}), \phi(\mathbf{x}')) \quad (6.19)$$

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{A} \mathbf{x}' \quad (6.20)$$

$$k(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}_a, \mathbf{x}'_a) + k_b(\mathbf{x}_b, \mathbf{x}'_b) \quad (6.21)$$

$$k(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}_a, \mathbf{x}'_a)k_b(\mathbf{x}_b, \mathbf{x}'_b) \quad (6.22)$$

where $c > 0$ is a constant, $f(\cdot)$ is any function, $q(\cdot)$ is a polynomial with nonnegative coefficients, $\phi(\mathbf{x})$ is a function from \mathbf{x} to \mathbb{R}^M , $k_3(\cdot, \cdot)$ is a valid kernel in \mathbb{R}^M , \mathbf{A} is a symmetric positive semidefinite matrix, \mathbf{x}_a and \mathbf{x}_b are variables (not necessarily disjoint) with $\mathbf{x} = (\mathbf{x}_a, \mathbf{x}_b)$, and k_a and k_b are valid kernel functions over their respective spaces.

How to Prove If a Kernel Is Valid?

1. Prove that there exists a ϕ such that
$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}'), \forall \mathbf{x}, \mathbf{x}'$$
2. Prove that the Gram matrix K is PSD and use Mercer's Theorem
 - Note: PSD + PSD = PSD and $c \times \text{PSD} = \text{PSD}$ for $c \geq 0$
 - Also useful to prove if a kernel is invalid; provide a counterexample showing that the Gram matrix K is not PSD
3. Use the axioms provided in previous slides
 - But **not for assignments & exams**; you need to prove them before using them.

Commonly Used Kernels

- Simple polynomial kernel (degree of 2)

$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z})^2$$

- Generalized polynomial kernel (degree of M)

$$k(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z} + c)^p, c > 0$$

- Gaussian kernel

$$k(\mathbf{x}, \mathbf{z}) = \exp\left(-\frac{\|\mathbf{x} - \mathbf{z}\|^2}{2\sigma^2}\right)$$

- Cf. Radial basis function (RBF) kernel is essentially the same, but has different parametrization.

$$k(\mathbf{x}, \mathbf{z}) = \exp(-\gamma \|\mathbf{x} - \mathbf{z}\|^2)$$

Gaussian Kernel

- Not actually Gaussian pdf
- Translation invariant; depends only on distance between points, such that it can be expressed as
$$k(\mathbf{x}, \mathbf{z}) = k(\mathbf{x} - \mathbf{z})$$
- Corresponds to a mapping to an infinite-dimensional space!

(Nadaraya-Watson) Kernel Regression

Kernel Regression

- Intuition: If a test data is close to a training data, then they would have a similar target.
 - Take a kernel as the similarity measure
 - Take the weighted average of targets of training data, where the weight is the kernel value
 - No training is required, similar to K-nearest neighbors

- Kernel regression output

$$h(\mathbf{x}) = \frac{\sum_n k(\mathbf{x}, \mathbf{x}^{(n)}) y^{(n)}}{\sum_n k(\mathbf{x}, \mathbf{x}^{(n)})}$$

- Note that this is different from kernel ridge regression.

Kernel Regression vs. Classification

- Similar to K-nearest neighbors, we can reformulate it into classification.

- Kernel regression

$$h(\mathbf{x}) = \frac{\sum_n k(\mathbf{x}, \mathbf{x}^{(n)}) y^{(n)}}{\sum_n k(\mathbf{x}, \mathbf{x}^{(n)})}$$

- Kernel classification (for $y \in \{-1, +1\}$)

$$h(\mathbf{x}) = \text{sign} \left(\sum_n k(\mathbf{x}, \mathbf{x}^{(n)}) y^{(n)} \right)$$

vs. Locally-Weighted Linear Regression

- Kernel regression
 - No training
 - Output: $h(\mathbf{x}) = \frac{\sum_n k(\mathbf{x}, \mathbf{x}^{(n)}) y^{(n)}}{\sum_n k(\mathbf{x}, \mathbf{x}^{(n)})}$
- Locally-weighted linear regression
 - Find \mathbf{w} to minimize $J(\mathbf{w}) = \sum_n r^{(n)} (\mathbf{w}^T \mathbf{x}^{(n)} - y^{(n)})^2$
 - Output: $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}^{(n)}$
- Common choice for both $k(\mathbf{x}, \mathbf{x}^{(n)})$ and $r^{(n)}$ is Gaussian kernel:

$$\exp\left(-\frac{\|\mathbf{x} - \mathbf{x}^{(n)}\|^2}{2\sigma^2}\right)$$

vs. Locally-Weighted Linear Regression

- Similar: Instance-based learning
 - Only observations (training data) close to the query point are considered (highly weighted) for regression.
 - Kernel determines how much weights to training data by computing similarity to the query (test data).
 - Free to choose any kernel
 - Both can suffer when the input dimension is high.
- Dissimilar:
 - Kernel regression does not perform training.
 - In general, kernel regression is faster but less accurate.

Next: Support Vector Machines